

Trading the SP500 with Reinforcement Learning (1)

April 22, 2020

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from bokeh.io import show, output_notebook
from bokeh.layouts import column
from bokeh.plotting import figure, output_file, show
from bokeh.layouts import layout, row
from bokeh.models import ColumnDataSource, HoverTool, Span, BoxAnnotation
from datetime import datetime
output_notebook()

[2]: hover_price = HoverTool(
    tooltips=[
        ('date', '@date{%F}'),
        ('price', '$@{price}{0.2}'),],
    formatters={
        '@date': 'datetime',
        '@{price}': "numeral",
    },
    mode='vline')

[3]: hover_change = HoverTool(
    tooltips=[
        ('date', '@date{%F}'),
        ('change', '@{change}{.0000}'),],
    formatters={
        '@date': 'datetime',
        '@{change}': "numeral",
    },
    mode='vline')

[4]: sp500 = pd.read_csv("../data/GSPC_2011.csv", parse_dates=["Date"])
sp500["price_change"] = sp500["Close"].pct_change(1)+1
sp500 = sp500.dropna()

price = ColumnDataSource(data={"date":sp500["Date"], "price":sp500['Close']})
```

```

p = figure(title="SP500 Price", plot_height=300, plot_width=900, tools="xpan",
→toolbar_location=None,
           x_axis_type="datetime",background_fill_color="#efefef",
→x_range=(sp500["Date"].iloc[0], sp500["Date"].iloc[-1]))
p.line('date', 'price', source=price)
p.yaxis.axis_label = 'Price'
p.add_tools(hover_price)

price_change = ColumnDataSource(data={"date":sp500["Date"], "change":
→sp500['price_change']})
c = figure(title="SP500 Daily Price Change", plot_height=300, plot_width=900,
→tools="xpan", toolbar_location=None,
           x_axis_type="datetime",background_fill_color="#efefef",
→x_range=(sp500["Date"].iloc[0], sp500["Date"].iloc[-1]))

c.line('date', 'change', source=price_change)
c.yaxis.axis_label = 'Change'
c.add_tools(hover_change)

```

8

9 1 Trading the SP500 with Reinforcement Learning

10 1.0.1 Abstract

11 Can a machine learning algorithm outperform the stock market? This was the question we sought
12 to answer. To do so we trained a deep q-learning algorithm with end of day prices of the S&P 500
13 in 2011 to recommend favorable trades. Our agent managed to beat the market even though 2011
14 was a very volatile year for the stock market. Nevertheless the agent is not ready to be used in
15 any real life scenario as it is way to volatile. More indicator and longer time frames are needed to
16 train and test this model until it can be used.

17 1.0.2 Introduction

18 The stock market is a game of risk and uncertainty. Every day thousands of people trade with
19 trillions of dollar to secure profits for themselves and their clients. How will the market move
20 next? That is the question everyone tries to answer and the lucky few that predict this consistently
21 can make billions of dollar. This game is played with intelligence and huge amounts of data but
22 unfortunately also a lot of emotions.

23 Algorithms are build to process huge amounts of data, they are always focussed on their task, can
24 act in milliseconds and dont have any emotions that could interfere. They are the perfect trader.

25 With recent advances in machine learning it seems only logical to ask if a machine learning algo-
26 rithm fueled by tons of data could outperform human trader. So we started this project with the
27 goal to build a **machine learning algorithm that could outperform the stock market**.

28 The goal here was two fold. On one hand we obviously want to teach an agent to recommend prof-

itable trades but even more important than that was to learn how machine learning and especially reinforcement learning works in general.

The plan was to create a deep q-learning algorithm that would learn from historical timeseries data of the S&P 500 and recommend trades that, if executed would generate more profits than the S&P 500 in the same time.

1.0.3 Materials and Methods

The S&P 500 is one of the most common indices containing the 500 biggest US-companies by market capitalization. We used daily closing prices (the last reported price of any trading day) from YahooFinance. There are more specialized data-provider out there with even better data but the S&P 500 is well reported and Yahoo is more than precise enough for our experiments with reinforcement learning.

The agent is trained using a Q-learning approach. We used a neural network and trained the algorithm with small batches of ten random successive trading days. Each time the algorithm chose an action that made a profit it got an reward and each time it lost money it suffered a punishment. This way the algorithm learns which actions are favorable and which are not. We used a Sequential model with three dense hidden layers which went through 300 training episodes. The agent gets ten successive data points and learns to predict the ideal action for the coming trading day. It can choose between doing nothing ("sit"), buying more ("buy") or sell ("sell"). If the agent chooses to buy or sell it can only buy/sell one unit and it can obviously only sell what it already owns. The model only chooses action it has no concept of a budget or trading costs.

```
model = Sequential()  
model.add(Dense(units=64, input_dim=self.state_size, activation="relu"))  
model.add(Dense(units=32, activation="relu"))  
model.add(Dense(units=8, activation="relu"))  
model.add(Dense(self.action_size, activation="linear"))  
model.compile(loss="mse", optimizer=Adam(lr=0.001))
```

We started to train the agent on daily data of the S&P 500 going from 2001 to 2011 but this proved simply to time intensive which is why we chose to only use data from the year 2011 for now. We chose 2011 cause we wanted to train the agent on a more complex data set, which also includes vast drawdowns instead of only a steady uptrend (which often leads to a simple buy and hold).

There is not much data processing happening on the agent side. We normalized the closing prices and broke them into chunks of 10 data points.

1.0.4 Results

Lets now speak about the results. In order to understand the underlying data on which the algorithm got trained we will go into a bit of history and show the major events happening in 2011 and how it impacted the S&P 500.

February 15th - The first Libyan Civil War starts. Uncertainty about libyan oil production impacts the stock market.

```
[5]: start_libyan_civil_war = Span(location=datetime(2011,2,15), dimension='height',
    ↳line_color='orange', line_width=1)
    p.add_layout(start_libyan_civil_war)
```

67

68 March 11th - Earthquake of strength 9.0 hits Japan. A Tsunami kills thousands of people and leads
 69 to a catastrophe in the nuclear reactor in Fukushima. Japans economy is in a shock which has
 70 ripple effects through other markets.

```
[6]: japan_earthquake = Span(location=datetime(2011,3,11), dimension='height',
    ↳line_color='purple', line_width=1)
    p.add_layout(japan_earthquake)
```

71

72 March 15th - Exactly one month after the start of the libyan civil war the president of Bahrai
 73 declares a state of emergency in response to continous protests and mobilizes the army against
 74 them. In Syria the civil war starts. Fear of oil- and energy shortages in the US increase with
 75 conflicts in Libya, Bahrain and Syria and civil unrest in other arabic countries.

```
[7]: bahrain_syria_conflicts = Span(location=datetime(2011,3,15), dimension='height',
    ↳line_color='coral', line_width=1)
    p.add_layout(bahrain_syria_conflicts)
```

76

77 June - Debates about the second bailout for greece lead to more and more uncertainty in the mar-
 78 kets about the future of greece, europe and the euro.

```
[8]: greek_baylout_debates = BoxAnnotation(left=datetime(2011,5,31),
    ↳right=datetime(2011,6,29), fill_color='blue', fill_alpha=0.1)
    p.add_layout(greek_baylout_debates)
```

79

80 June 29th - The third credit package for greece is confirmed. With it are huge austerity measures
 81 coming for greece and the greek public.

```
[9]: credit_greece = Span(location=datetime(2011,6,29), dimension='height',
    ↳line_color='green', line_width=1)
    p.add_layout(credit_greece)
```

82

83 August 5th - High debt and major difficulties regarding the US-Budget cause the US credit rating
 84 to drop from AAA to AA+. This is the first time since 1941 that the US doesn't count as AAA.

```
[10]: us_credit_downgrade = Span(location=datetime(2011,8,5), dimension='height',
    ↳line_color='red', line_width=1)
    p.add_layout(us_credit_downgrade)
```

85

86 August till December - The US stock market is in a major crisis fueled by the european debt crisis
 87 and fear about its effect on the us economy. Occupy Wall Street starts their protests in september.
 88 The rest of the the year is marked by uncertainty, doubt and fear. Each uptrend gets wiped out
 89 quickly and drastically. The market can recover a bit to the end of the year.

90 September 19th - Italian credit gets downgrade from A+ to A. This is a major hit for many banks
 91 after they lend billions to stabilize the country

```
[11]: italian_credit_downgrade = Span(location=datetime(2011,9,19),
    ↳ dimension='height', line_color='red', line_width=1)
p.add_layout(italian_credit_downgrade)
```

92

93 [1]

```
[12]: show(column(p, c))
```

94

95 **Figure 1: The SP500 Price** shows the daily closing price of the S&P 500 in the year 2011. Colored
 96 lines or areas show historic events.

97 **Figure 2: The SP500 Daily Price Change** tells us the daily change of price in percent.

98 Now that we know what happend in 2011 and how the market reacted, lets look at how our agent
 99 performed in this environment. We will first look into the actions it chose. Following this we will
 100 dissect the chosen trades and analyze the performance of it. We will conclude our results with an
 101 analysis of the viability of those recommendations.

```
[13]: dqn = pd.read_csv("./data/dqn_model.csv", parse_dates=["Date"])
dqn["Strategy"] = dqn["Profits"].cumsum()
dqn["Return_perc"] = dqn["Profits"]/dqn["Price"] + 1

buys = dqn[dqn["Action"] == "Buy"][["Date", "Price"]]
sells = dqn[dqn["Action"] == "Sell"][["Date", "Price"]]
```

102

```
[14]: p2 = figure(title="SP500 DQN Trading Agent", plot_height=300, plot_width=900,
    ↳ toolbar_location=None,
        x_axis_type="datetime", background_fill_color="#efefef",
    ↳ x_range=(dqn["Date"].iloc[0], dqn["Date"].iloc[-1]))
p2.line('date', 'price', source=price)
p2.yaxis.axis_label = 'Price'
p2.circle(buys["Date"], buys["Price"], size=5, color="green")
p2.circle(sells["Date"], sells["Price"], size=5, color="red")
show(p2)
```

103

104 **Figure 3:** Here we can see which action the algorithm chose. Green dots are buys and red dots
 105 show sells.

106 This looks overall quite reasonable. Not too much sticks out. Sure, here and there are some sell
 107 signals at local minima and buys at local maxima but it does not seem to extreme. So lets look a
 108 bit more into it. What did the agent actually do?

109 The agent sits a bit more than half ot the time just idle. Buys and Sells are even which makes sense.
 110 We did not sell more than we had and at the end of the year all our positions are closed.

```
[15]: dqn['Action'].value_counts()
```

111

```
[15]: Sit      130
      Sell     61
      Buy      61
```

112


Name: Action, dtype: int64

113

```
[16]: fig, (ax1, ax2) = plt.subplots(1,2, figsize=(10,4))
fig.suptitle('Trade Returns', size=16)
ax1.plot(dqn["Date"],dqn["Return_perc"]-1)
ax1.set_title("Returns over time")
ax1.set_xlabel("Month")
ax1.set_ylabel("Return Percentage")
ax1.tick_params(labelrotation=45)

ax2.hist(dqn[dqn["Return_perc"] != 1.0]["Return_perc"]-1, bins=30)
ax2.axvline((dqn[dqn["Action"] == "Sell"]["Return_perc"]-1).median(),
            color="red", label="median")
ax2.axvline((dqn[dqn["Action"] == "Sell"]["Return_perc"]-1).mean(),
            color="yellow", label="mean")
ax2.set_xlabel("Return Percentage")
ax2.set_ylabel("Occurences")
ax2.tick_params(labelrotation=45)
ax2.set_title("Return Distribution")
plt.show()
```

114



115

116

117

118 **Figure 4: Returns over time** shows the return percentage of each trade and when they occurred.


119 **Figure 5: Return Distribution** shows the distribution of returns per trade.

120 As we can see in Figure 4 and 5 most trades only have a marginal impact. The whole Strategy is
121 only really impacted by a few dramatic trades. Our median is pretty close to zero and the average
122 return is only 0.28% (red and yellow line in Figure 5).

123 Our best trade scores 7.25% but the worst trade costs -8.5%. Figure 1 shows nicely that a majority
124 of time is relatively calm. We are either not trading or just getting minor profits/losses. All these

125 average trades are dwarfed by our few big movers.
126 But how does this compare to buying and holding the S&P 500 in 2011?

```
[17]: sp500["price_change"].cumprod().plot()  
      dqn["Return_perc"].cumprod().plot()  
      plt.show()
```



128

129

130

131 **Figure 6:** This plot shows the returns in percent of the S&P 500 versus our algorithm. The orange
132 line is our algorithm and the blue line shows the S&P 500.

133 Actually surprisingly good. If we would have followed the recommendations of our agent we
134 would have made 15.7% instead of losing 1.2% in the S&P 500. Sounds great doesn't it?

135 Not so fast. While at the end of the year we would have potentially made more money than by
136 following the Index it is very unrealistic that anyone would have gotten to that place. After seeing
137 the agent recommending bad trade after bad trade and being permanently worse performing
138 than the index, only the most stubborn trader could still follow these recommendations. At the
139 lowest point the S&P 500 is down 13.6% against its initial price in January. Our agent though is
140 down 27.8% on its worst point.

```
[18]: print(f'SP500 Worst: {round(sp500["price_change"].cumprod().min()-1,3)}')  
      print(f'Agent Worst: {round(dqn["Return_perc"].cumprod().min()-1,3)}')
```

141

```
SP500 Worst: -0.136  
Agent Worst: -0.217
```

142 The S&P 500 is also very limited in its upside with only a maximum gain of 7.2% while our agent
143 delivers 26.1%.

```
[19]: print(f'SP500 Best: {round(sp500["price_change"].cumprod().max()-1,3)}')  
      print(f'Agent Best: {round(dqn["Return_perc"].cumprod().max()-1,3)}')
```

SP500 Best: 0.072

Agent Best: 0.261

And that's how we come to one of the biggest problems with the algorithm. It's way too volatile, even more so than the S&P 500 in a very volatile year. From the lowest point to its highest the agent more than doubles its return. But even the large trades in between are too volatile.

It should be clear by now that the agent in its current form doesn't give recommendations that could be applied in any real life scenario.

1.0.5 Discussion

2011 was a turbulent year, the beginning of the Arab Spring, an oil crisis, the European debt crisis and it was the first time since 1941 that the US credit score dropped below AAA. The S&P 500 closes with a loss of 1.2% from its initial price at the start of the year. Nonetheless our agent made 15.7% in this environment.

These returns are unfortunately not actionable. Most of the year our algorithm performs worse than the S&P 500 and only to the end of the year we gain all our profits in a few great trades. Our agent's recommendations result in not only a strategy that is way too volatile but even those huge gains could simply be a result of luck.

It is hereby important to put this into context. We only used data from 2011, 300 training episodes and only closing prices. No indicator or other complementary timeseries. I am therefore very optimistic that this result could be improved by quite a bit.

There are big Hedge Funds that rely exclusively on machine learning algorithms to generate their profits. It makes sense, machines are better at processing huge amounts of data and don't have any emotions that could interfere. But it also makes sense that these successful algorithms are not easy to implement. If that would be the case none of those firms would need a fast staff of researchers and developers.

We can now answer our question if a reinforcement learning algorithm can beat the market. Yes it can, but it still takes a lot of work to make the algorithm reliable and consistent to the degree that it can be actually used in the stock market.

Leon Niesler, 13.04.2020

1.1 References

[1] <https://de.wikipedia.org/wiki/2011>

[]:
173

[]:
174

[]:
175